

1. Download the dataset: Dataset

2. Load the dataset

```
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import seaborn as sbn
%matplotlib inline
file=pd.read_csv("abalone.csv")
df=pd.DataFrame(file)
df.head()
Sex    Length Diameter      Height Whole weight  Shucked weight  Viscera weight
Shell weight  Rings
0      M      0.455 0.365 0.095 0.5140 0.2245 0.1010 0.150 15
1      M      0.350 0.265 0.090 0.2255 0.0995 0.0485 0.070 70
2      F      .530 0.420 0.135 0.6770 0.2565 0.1415 0.210 90.
3      M      440 0.365 0.125 0.5160 0.2155 0.1140 0.155 100.
4      I      330 0.255 0.080 0.2050 0.0895 0.0395 0.055 7
df['age'] = df['Rings']+1.5
df = df.drop('Rings', axis = 1)
```

3. Perform Below Visualizations.


- Univariate Analysis
- Bi - Variate Analysis
- Multi - Variate Analysis

#Univariate Analysis

```
sbn.boxplot(df.Length)
```

/usr/local/lib/python3.7/dist-packages/seaborn/_decorators.py:43: FutureWarning: Pass the following variable as a keyword arg: x. From version 0.12, the only valid positional argument will be `data`, and passing other arguments without an explicit keyword will result in an error or misinterpretation.


FutureWarning



the data is significantly imbalanced


#Bi-Variate Analysis

```
sbn.barplot(x=df.Height,y=df.Diameter)
```



#Multi-Variate Analysis

```
sbn.pairplot(df)
```



4. Perform descriptive statistics on the dataset.

df.info()

RangeIndex: 4177 entries, 0 to 4176

Data columns (total 9 columns):

Column Non-Null Count Dtype

```

---
0 Sex          4177 non-null object
1 Length       4177 non-null float64
2 Diameter     4177 non-null float64
3 Height       4177 non-null float64
4 Whole weight 4177 non-null float64
5 Shucked weight 4177 non-null float64
6 Viscera weight 4177 non-null float64
7 Shell weight 4177 non-null float64
8 age          4177 non-null float64

```

dtypes: float64(8), object(1)

memory usage: 293.8+ KB

df.describe()

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight
count	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000	4177.000000
mean	0.523992	0.238831	0.407881	0.139516	0.828742	0.359367	0.180594
std	0.120093	0.139203	0.099240	0.041827	0.490389	0.221963	0.109614
min	0.075000	0.001500	0.055000	0.000000	0.002000	0.001000	0.000500
25%	0.450000	0.130000	0.350000	0.115000	0.441500	0.186000	0.093500
75%	0.600000	0.250000	0.500000	0.150000	0.600000	0.300000	0.150000

50%	0.545000	0.425000	0.140000	0.799500	0.336000	0.171000
	0.234000	10.500000				
75%	0.615000	0.480000	0.165000	1.153000	0.502000	0.253000
	0.329000	12.500000				
max	0.815000	0.650000	1.130000	2.825500	1.488000	0.760000
	1.005000	30.500000				

5. Handle the Missing values.

```
df.isna().sum()
```

```
Sex      0
```

```
Length   0
```

```
Diameter 0
```

```
Height   0
```

```
Whole weight 0
```

```
Shucked weight 0
```

```
Viscera weight 0
```

```
Shell weight 0
```

```
age      0
```

```
dtype: int64
```

there is no missing values in dataset

for i in df:

```
    if df[i].dtype=='object' or df[i].dtype=='category':
```

```
        print("unique of "+i+" is "+str(len(set(df[i])))+" they are "+str(set(df[i])))
```

unique of Sex is 3 they are {'F', 'M', 'I'}

6. Find the outliers and replace the outliers

Checking for outliers

```
#Data Preprocessing
```

```
#Outlier handling
```


```
df = pd.get_dummies(df)
```

```
dummy_df = df
```

```
var = 'Viscera weight'
```

```
plt.scatter(x = df[var], y = df['age'])
```


```
plt.grid(True)
```



```
var = 'Shell weight'
```

```
plt.scatter(x = df[var], y = df['age'])
```


```
plt.grid(True)
```




```
var = 'Shucked weight'
```

```
plt.scatter(x = df[var], y = df['age'])
```


```
plt.grid(True)
```



```
var = 'Whole weight'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```




```
var = 'Diameter'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```




```
var = 'Height'
plt.scatter(x = df[var], y = df['age'])
```

```
plt.grid(True)
```



```
var = 'Length'
plt.scatter(x = df[var], y = df['age'])
plt.grid(True)
```



Removing outliers

```
df.drop(df[(df['Viscera weight'] > 0.5) &
           (df['age'] < 20)].index, inplace = True)
df.drop(df[(df['Viscera weight'] < 0.5) & (
df['age'] > 25)].index, inplace = True)
df.drop(df[(df['Shell weight'] > 0.6) & (df['age'] < 25)].index, inplace = True)
df.drop(df[(df['Shell weight'] < 0.8) & (df['age'] > 25)].index, inplace = True)
df.drop(df[(df['Shucked weight'] >= 1) & (df['age'] < 20)].index, inplace = True)
df.drop(df[(df['Viscera weight'] < 1) & (df['age'] > 20)].index, inplace = True)
df.drop(df[(df['Diameter'] < 0.1) & (df['age'] < 5)].index, inplace = True)
df.drop(df[(df['Diameter'] < 0.6) & (df['age'] > 25)].index, inplace = True)
df.drop(df[(df['Diameter'] >= 0.6) & (df['age'] < 25)].index, inplace = True)
df.drop(df[(df['Height'] > 0.4) & (df['age'] < 15)].index, inplace = True)
df.drop(df[(df['Height'] < 0.4) & (df['age'] > 25)].index, inplace = True)
df.drop(df[(df['Length'] < 0.1) & (df['age'] < 5)].index, inplace = True)
df.drop(df[(df['Length'] < 0.8) & (df['age'] > 25)].index, inplace = True)
```

```
df.drop(df[(df['Length']>=0.8) & (df['age'] < 25)].index, inplace = True)
```

7. Check for Categorical columns and perform encoding.

```
from sklearn.preprocessing import LabelEncoder
```

```
encoder=LabelEncoder()
```

```
for i in df:
```

```
    if df[i].dtype=='object' or df[i].dtype=='category':
```

```
        df[i]=encoder.fit_transform(df[i])
```

8. Split the data into dependent and independent variables.

```
x=df.iloc[:, :-1]
```

```
x.head()
```

	Length	Diameter	Height	Whole weight	Shucked weight	Viscera weight	Shell weight	age	Sex_F	Sex_M
0	0.455	0.365	0.095	0.5140	0.2245	0.1010	0.150	16.5	0	0
1	0.350	0.265	0.090	0.2255	0.0995	0.0485	0.070	8.5	0	0
2	0.530	0.420	0.135	0.6770	0.2565	0.1415	0.210	10.5	1	0
3	0.440	0.365	0.125	0.5160	0.2155	0.1140	0.155	11.5	0	0
4	0.330	0.255	0.080	0.2050	0.0895	0.0395	0.055	8.5	0	1

```
y=df.iloc[:, -1]
```

```
y.head()
```

```
0    1
```

```
1    1
```

```
2    0
```

```
3    1
```

```
4    0
```

```
Name: Sex_M, dtype: uint8
```

9. Scale the independent variables

```
from sklearn.preprocessing import StandardScaler
```

```
scaler=StandardScaler()
```

```
x=scaler.fit_transform(x)
```

```
x
```

```
array([[ -0.53701309, -0.39082366, -1.12698145, ...,  1.9433912 ,  
        -0.66579302, -0.70803622],  
       [ -1.42965864, -1.4205279 , -1.26123393, ..., -0.95032771,  
        -0.66579302, -0.70803622],  
       [  0.10059087,  0.17551367, -0.05296168, ..., -0.22689798,  
        1.50196828, -0.70803622],  
       ...,  
       [  0.6956879 ,  0.741851  ,  1.82657293, ..., -0.22689798,  
        -0.66579302, -0.70803622],  
       [  0.90822255,  0.84482142,  0.34979574, ...,  0.13481688,  
        1.50196828, -0.70803622],  
       [  1.63084038,  1.56561439,  1.55806799, ...,  0.85824661,  
        -0.66579302, -0.70803622]])
```

10. Split the data into training and testing

```
from sklearn.model_selection import train_test_split
```

```
x_train,x_test,y_train,y_test=train_test_split(x,y,test_size=0.33)
```

```
x_train.shape
```

```
(2676, 10)
```

```
x_test.shape
(1319, 10)
y_train.shape
(2676,)
y_test.shape
(1319,)
```

MODEL

Linear regression

```
from sklearn.linear_model import LinearRegression
lm = LinearRegression()
lm.fit(x_train, y_train)
LinearRegression()
y_train_pred = lm.predict(x_train)
y_test_pred = lm.predict(x_test)
from sklearn.metrics import mean_absolute_error, mean_squared_error
s = mean_squared_error(y_train, y_train_pred)
print('Mean Squared Error of training set :%.2f%s')
```

```
p = mean_squared_error(y_test, y_test_pred)
print('Mean Squared Error of testing set :%.2f%p)
Mean Squared Error of training set :0.000000
Mean Squared Error of testing set :0.000000
Note: The Lower the Mean Squared Error,better the forecast.
```

```
from sklearn.metrics import r2_score
s = r2_score(y_train, y_train_pred)
print('R2 Score of training set:%.2f%s')
```

```
p = r2_score(y_test, y_test_pred)
print('R2 Score of testing set:%.2f%p)
R2 Score of training set:1.00
R2 Score of testing set:1.00
Note: The ideal value of R-square is 1.
```

The closer the value of R-square to 1,better is the model fitted.